

LAB Logbook

Github repository link:

https://github.com/Salimboevm/ML_Finance

Lab 1

For the lab work in week 1, students were asked to create a vector using np.arange method and do some changes on the vector to practice Numpy and Python.

My SID is 1919019, and because of that created 19 elements(last two digits). Then, transformed this vector into a 2 dimensional array with 1 row using the reshape() method. Then, used NumPy's empty_like() method and slicing to create an independent array and save the values of the vector to that independent array. And finally, checked the shape attribute values of both arrays and printed all results at the end of each step.

Results:

```
[302]: vector = np.arange(19)
print(vector)
```

```
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18]
```

```
[306]: vector = vector.reshape(19,1)
print(vector)
```

```
[[ 0]
 [ 1]
 [ 2]
 [ 3]
 [ 4]
 [ 5]
 [ 6]
 [ 7]
 [ 8]
 [ 9]
 [10]
 [11]
 [12]
 [13]
 [14]
 [15]
 [16]
 [17]
 [18]]
```

```
[308]: new_array_2d = np.empty_like(vector)
new_array_2d[:, :] = vector
print(new_array_2d)
```

```
[[ 0]
 [ 1]
 [ 2]
 [ 3]
 [ 4]
 [ 5]
 [ 6]
 [ 7]
 [ 8]
 [ 9]
 [10]
 [11]
 [12]
 [13]
 [14]
 [15]
 [16]
 [17]
 [18]]
```

```
[310]: print(vector.shape)
print(new_array_2d.shape)
```

```
(19, 1)
(19, 1)
```

Lab 2

Lab 2 in Week 2, Pandas and its main functions were learnt. Requirements were using "adult_data_mini.csv" dataset and performing several operations.

Lab logbook requirement, n was determined as 9 (n=9) because of my student ID (last digit). Then data was grouped by "relationship" and "hours-per-week". Followed by "hours-per-week" column values were reduced by n=9. At this step, the function change_data(x) was created and used. To apply this function to the dataset, the apply() method was used and the original DataFrame was updated. Lastly, grouping by "relationship" and reduced "hours-per-week" was performed again.

Results:

```
Student ID: 1919019

[223]: group_by_hours = data.groupby(['relationship', 'hours-per-week'])
group_by_hours.size()

[223]:   relationship  hours-per-week
      Husband        13.0          1
                  40.0          2
                  45.0          1
                  80.0          1
    Not-in-family    16.0          1
                  40.0          2
                  50.0          2
  Own-child        30.0          1
    Wife           40.0          2
dtype: int64

[227]: def change_data(x):
        return x - 9

data['hours-per-week'] = data['hours-per-week'].apply(change_data)
data

[227]:   age  workclass  fnlwgt  education  education-num  marital-status  occupation  relationship  race  sex  capital-gain  capital-loss  hours-per-week  native-country  Answer  IsHomeles
      0  State-gov  77516.0  Bachelors       13.0  Never-married  Adm-clerical  Not-in-family  White  Male    2174.0     NaN    31.0  United-States  <=50K  False
      1  Self-emp-not-inc  83311.0  Bachelors       13.0  Married-civ-spouse  Exec-managerial  Husband  White  Male      0.0     0.0     4.0  United-States  <=50K  False
      2    Private  215646.0   HS-grad        9.0  Divorced  Handlers-cleaners  Not-in-family  White  Male      0.0     NaN    31.0  United-States  <=50K  False
      3    Private  234721.0      11th        7.0  Married-civ-spouse  Handlers-cleaners  Husband  Black  Male      0.0     NaN    31.0  United-States  <=50K  False
      4    Private  338409.0  Bachelors       13.0  Married-civ-spouse  Prof-specialty  Wife  Black  Female     0.0     NaN    31.0     Cuba  <=50K  False
      5    Private  284582.0    Masters       14.0  Married-civ-spouse  Exec-managerial  Wife  White  Female     0.0     NaN    31.0  United-States  <=50K  False
      6    Private  160187.0      9th        5.0  Married-spouse-absent  Other-service  Not-in-family  Black  Female     0.0     0.0     7.0   Jamaica  <=50K  False
      7  Self-emp-not-inc  209642.0   HS-grad       9.0  Married-civ-spouse  Exec-managerial  Husband  White  Male      0.0     0.0    36.0  United-States  >50K  False
      8    Private  45781.0     Masters       14.0  Never-married  Prof-specialty  Not-in-family  White  Female  14084.0     NaN    41.0  United-States  >50K  False
```

8	31	Private	45781.0	Masters	14.0	Never-married	Prof-specialty	Not-in-family	White	Female	14084.0	NaN	41.0	United-States	>50K	Fals
10	37	Private	280464.0	Some-college	10.0	Married-civ-spouse	Exec-managerial	Husband	Black	Male	0.0	NaN	71.0	United-States	>50K	Fals
12	23	Private	122272.0	Bachelors	13.0	Never-married	Adm-clerical	Own-child	White	Female	0.0	NaN	21.0	United-States	<=50K	Fals
13	32	Private	205019.0	Assoc-acdm	12.0	Never-married	Sales	Not-in-family	Black	Male	0.0	NaN	41.0	United-States	<=50K	Fals
14	40	Private	121772.0	Assoc-voc	11.0	Married-civ-spouse	Craft-repair	Husband	Asian-Pac-Islander	Male	0.0	NaN	31.0	?	>50K	Fals
15	25	Private	NaN	Some-college	NaN	NaN	NaN	NaN	White	Male	0.0	NaN	NaN	NaN	NaN	Tru


```
[229]: group_by_reduced_hours = data.groupby(['relationship', 'hours-per-week'])
group_by_reduced_hours.size()
```

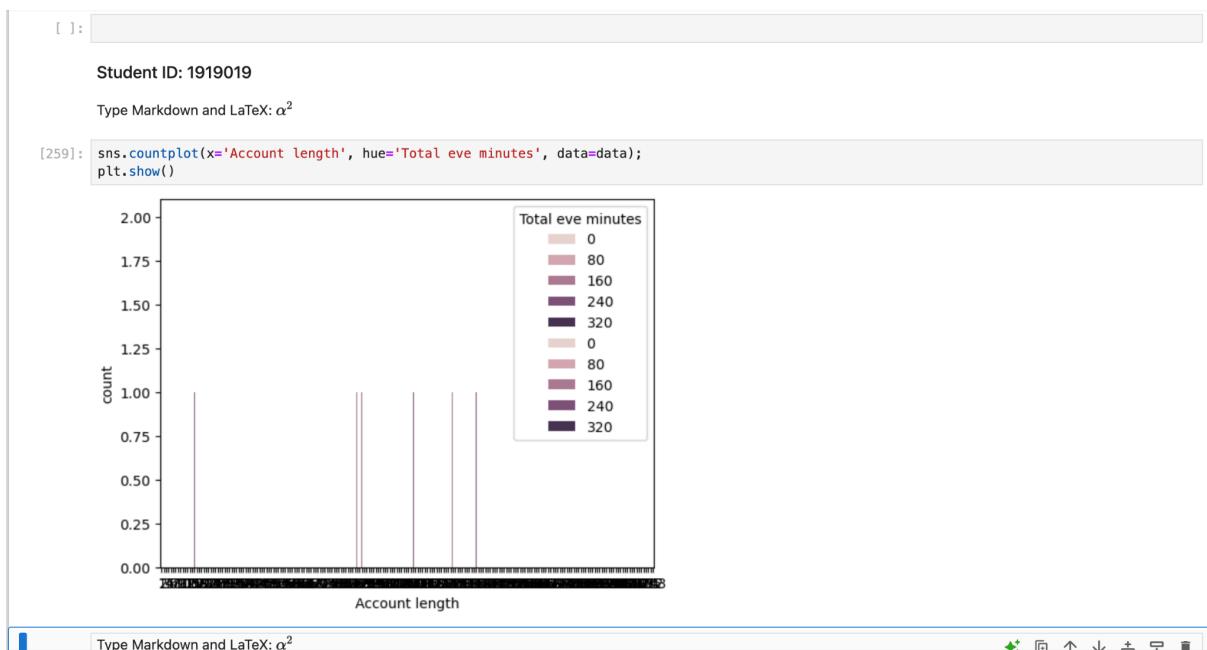
[229]:	relationship	hours-per-week	
	Husband	4.0	1
		31.0	2
		36.0	1
		71.0	1
	Not-in-family	7.0	1
		31.0	2
		41.0	2
	Own-child	21.0	1
	Wife	31.0	2
			dtype: int64

Lab 3

Lab 3 in Week 3, a bicolour features interaction diagram drawing was requested as a requirement. Because of my student ID, the selected columns are the 1st and 9th columns (based on last two digits: 19).

According to the diagram, the visualisation shows the interaction between these two features using a bicolour scheme. The Seaborn pairplot() function with a hue parameter was used to create the bicolour effect, allowing for clear visual distinction between different categories in the dataset.

Results:



Type Markdown and LaTeX: α^2



Lab 4

Lab 4 in Week 4, a Multilayer Perceptron (MLP) Neural Network was developed in the practical session. This was the first price forecast model, and the S&P-500 Index Prices dataset was used. For Lab 4, students were requested to create their own MLP model with two hidden layers.

Because of my student ID, the cells inside my hidden layers are 19 for the first hidden layer and 10 for the second hidden layer (approximately half of 19). There is also an output layer with one cell. After model creation, compiled the model and trained it with the same datasets from the practical session for 10 epochs. Finally, received the Mean Absolute Error (MAE) result.

Results:

```
[112]: model_2 = keras.Sequential([
    keras.layers.Dense(19, input_dim=500, activation=tf.nn.relu, kernel_initializer="normal"),
    keras.layers.Dense(10, activation='relu', kernel_initializer='normal'),
    #Output layer
    keras.layers.Dense(1)
])
print(model_2.summary())
Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 19)	9,519
dense_9 (Dense)	(None, 10)	200
dense_10 (Dense)	(None, 1)	11

```
Total params: 9,730 (38.01 KB)
Trainable params: 9,730 (38.01 KB)
Non-trainable params: 0 (0.00 B)
None
[114]: # because my SID is 1919019 and last three digits are 019, I can't use 019 instead I used 19
[116]: model_2.compile(optimizer="adam", loss="mse", metrics=["mae"])
[118]: history_2 = model_2.fit(X_train, y_train, batch_size=10, epochs=10, validation_split=0.2, verbose=1)
Epoch 1/10
2640/2640 - 2s 729us/step - loss: 6.2650e-04 - mae: 0.0142 - val_loss: 0.0049 - val_mae: 0.0639
Epoch 2/10
```

Model Summary:

- Layer 1 (Dense): Output Shape (None, 19), Parameters: 9,519
- Layer 2 (Dense): Output Shape (None, 10), Parameters: 200
- Layer 3 (Dense): Output Shape (None, 1), Parameters: 11
- Total parameters: 9,730

```
[118]: history_2 = model_2.fit(X_train, y_train, batch_size=10, epochs=10, validation_split=0.2, verbose=1)
Epoch 1/10
2640/2640    2s 729us/step - loss: 6.2650e-04 - mae: 0.0142 - val_loss: 0.0049 - val_mae: 0.0639
Epoch 2/10
2640/2640    1s 484us/step - loss: 1.0494e-04 - mae: 0.0079 - val_loss: 0.0015 - val_mae: 0.0332
Epoch 3/10
2640/2640    1s 380us/step - loss: 7.9832e-05 - mae: 0.0069 - val_loss: 0.0013 - val_mae: 0.0304
Epoch 4/10
2640/2640    1s 486us/step - loss: 6.8821e-05 - mae: 0.0064 - val_loss: 0.0012 - val_mae: 0.0287
Epoch 5/10
2640/2640    1s 432us/step - loss: 6.1618e-05 - mae: 0.0059 - val_loss: 0.0018 - val_mae: 0.0375
Epoch 6/10
2640/2640    2s 573us/step - loss: 5.2630e-05 - mae: 0.0055 - val_loss: 8.2489e-04 - val_mae: 0.0234
Epoch 7/10
2640/2640    1s 506us/step - loss: 4.9318e-05 - mae: 0.0053 - val_loss: 6.4474e-04 - val_mae: 0.0203
Epoch 8/10
2640/2640    1s 553us/step - loss: 5.0440e-05 - mae: 0.0054 - val_loss: 6.4466e-04 - val_mae: 0.0205
Epoch 9/10
2640/2640    1s 405us/step - loss: 4.8714e-05 - mae: 0.0052 - val_loss: 5.1350e-04 - val_mae: 0.0183
Epoch 10/10
2640/2640    1s 417us/step - loss: 4.6789e-05 - mae: 0.0051 - val_loss: 3.1803e-04 - val_mae: 0.0149

[120]: mse_2, mae_2 = model_2.evaluate(X_test, y_test, verbose=0)
print("Mean absolute error: %.5f" % mae_2)

Mean absolute error: 0.01643
```

Lab 5

Lab 5 in Week 5, a CNN code example was written in the practical session for price forecasting of Forex EUR/USD. For the lab logbook requirement, another CNN model creation was requested.

In this assignment, the convolutional core size (kernel_size) should be taken as 5, and batch_size should be 50. Additionally, the number of epochs should be determined according to Student IDs. Because of my student number, with Z=1 and Y=9, and after calculation it was 10. After compiling and training my model, I received the MAE value as a result.

Mean Absolute Error shows us the average of absolute differences between actual values and predicted values. The comparison between my CNN and the practical session CNN reveals how parameter differences (kernel_size, batch_size, and epochs) affect model performance. The CNN with lower MAE demonstrates better predictive capability.

Results:

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv1d_4 (Conv1D)	(None, 50, 50)	1,300
max_pooling1d_2 (MaxPooling1D)	(None, 7, 50)	0
conv1d_5 (Conv1D)	(None, 7, 100)	25,100
global_max_pooling1d_2 (GlobalMaxPooling1D)	(None, 100)	0
dense_4 (Dense)	(None, 25)	2,525
dense_5 (Dense)	(None, 2)	52

Total params: 28,977 (113.19 KB)

Trainable params: 28,977 (113.19 KB)

Non-trainable params: 0 (0.00 B)

None

```
[106]: model_2.compile(optimizer="adam", loss="mse", metrics=["mae"])
```

```
[108]: history_2 = model_2.fit(X_train, y_train, batch_size=50, epochs=10, validation_split=0.2, verbose=1)
```

```
Epoch 1/10
3520/3520 8s 2ms/step - loss: 0.0025 - mae: 0.0277 - val_loss: 9.1150e-04 - val_mae: 0.0197
Epoch 2/10
3520/3520 6s 2ms/step - loss: 7.5722e-04 - mae: 0.0188 - val_loss: 9.0916e-04 - val_mae: 0.0202
Epoch 3/10
3520/3520 6s 2ms/step - loss: 7.1741e-04 - mae: 0.0181 - val_loss: 9.6782e-04 - val_mae: 0.0214
Epoch 4/10
3520/3520 9s 2ms/step - loss: 7.1231e-04 - mae: 0.0180 - val_loss: 8.5019e-04 - val_mae: 0.0191
Epoch 5/10
3520/3520 9s 3ms/step - loss: 7.0645e-04 - mae: 0.0179 - val_loss: 8.4837e-04 - val_mae: 0.0188
```

```
[100]: ##### Student ID: 1919019
```

```
[102]: #Number of epochs: Z + Y = 10
#(Z = 1
#Y = 9)
```

```
[104]: model_2 = keras.Sequential([
    keras.layers.Conv1D(50, 5, padding='same', input_shape=(50,5), activation=tf.nn.relu, kernel_initializer="normal"),
    keras.layers.MaxPooling1D(7),
    keras.layers.Conv1D(100, 5, padding='same', activation=tf.nn.relu, kernel_initializer="normal"),
    keras.layers.GlobalMaxPooling1D(),
    keras.layers.Dense(25, activation=tf.nn.relu, kernel_initializer="normal"),
    keras.layers.Dense(2)
])
print(model_2.summary())
Model: "sequential_2"
```

```
[108]: history_2 = model_2.fit(X_train, y_train, batch_size=50, epochs=10, validation_split=0.2, verbose=1)
```

```
Epoch 1/10
3520/3520 8s 2ms/step - loss: 0.0025 - mae: 0.0277 - val_loss: 9.1150e-04 - val_mae: 0.0197
Epoch 2/10
3520/3520 6s 2ms/step - loss: 7.5722e-04 - mae: 0.0188 - val_loss: 9.0916e-04 - val_mae: 0.0202
Epoch 3/10
3520/3520 6s 2ms/step - loss: 7.1741e-04 - mae: 0.0181 - val_loss: 9.6782e-04 - val_mae: 0.0214
Epoch 4/10
3520/3520 9s 2ms/step - loss: 7.1231e-04 - mae: 0.0180 - val_loss: 8.5019e-04 - val_mae: 0.0191
Epoch 5/10
3520/3520 9s 3ms/step - loss: 7.0645e-04 - mae: 0.0179 - val_loss: 8.4837e-04 - val_mae: 0.0188
Epoch 6/10
3520/3520 8s 2ms/step - loss: 7.0028e-04 - mae: 0.0178 - val_loss: 9.6686e-04 - val_mae: 0.0213
Epoch 7/10
3520/3520 8s 2ms/step - loss: 6.9346e-04 - mae: 0.0176 - val_loss: 8.2854e-04 - val_mae: 0.0184
Epoch 8/10
3520/3520 7s 2ms/step - loss: 6.9649e-04 - mae: 0.0177 - val_loss: 8.2802e-04 - val_mae: 0.0185
Epoch 9/10
3520/3520 7s 2ms/step - loss: 6.9377e-04 - mae: 0.0177 - val_loss: 8.3583e-04 - val_mae: 0.0187
Epoch 10/10
3520/3520 7s 2ms/step - loss: 6.9154e-04 - mae: 0.0176 - val_loss: 8.3504e-04 - val_mae: 0.0187
```

```
[109]: mse, mae = model_2.evaluate(X_test, y_test, verbose=1)
print("Mean absolute error: %.5f" % mae)
```

```
936/936 1s 601us/step - loss: 0.0013 - mae: 0.0245
Mean absolute error: 0.02453
```

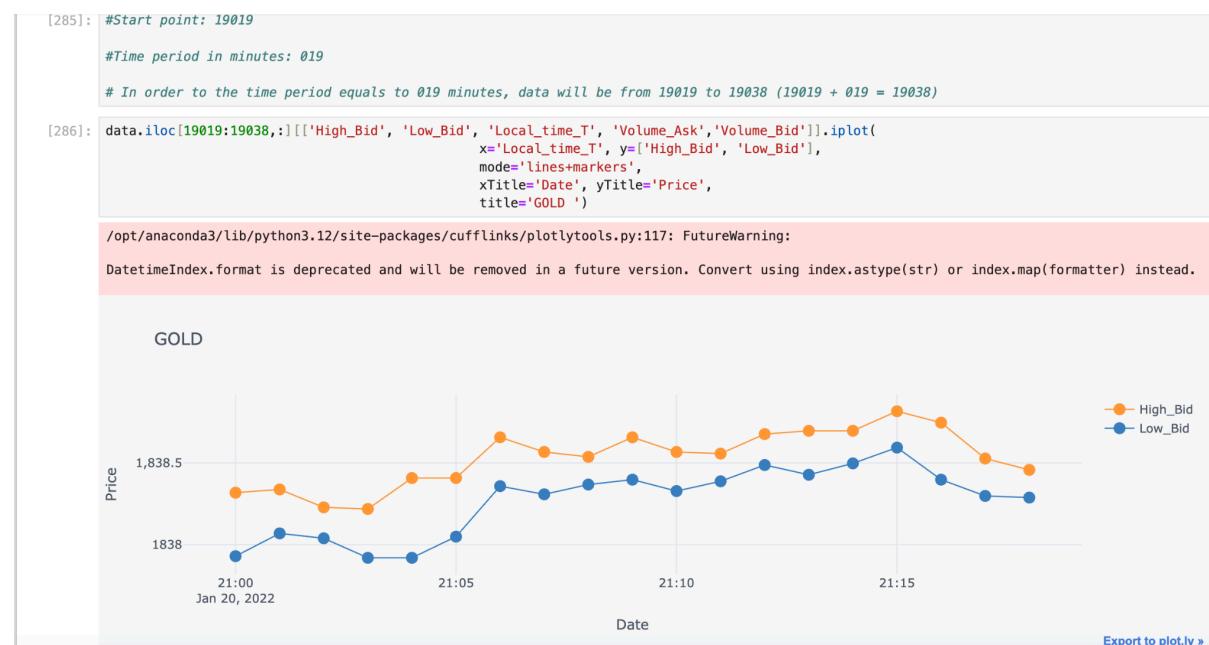
```
[ ]:
```

Lab 6

Lab 6 in Week 6, detailed code for data preprocessing was written. For the Lab Logbook requirement, students were asked to plot the price chart of a part of the whole dataset showing 'High_Bid' and 'Low_Bid' prices using the iplot() library.

Because my student ID is 1919019, the start point = 19019 and the time period in minutes = 19 (last three digits = 019 = 19). In order for the time period to equal 19 minutes, the data in the graph will be from index 19019 to 19038 (19019 + 19 = 19038).

Results:



Lab 7

Lab 7 in Week 7, LSTM and early stopping were explored. For the Lab Logbook requirement, a new LSTM model should be created with the following parameters:

- LSTM model parameter: $ZY + 10 = 19 + 10 = 29$
(Because of my Student ID, $ZY = 19$)
- Epochs = 10
- Patience = 3

As a result, I obtained the MAE and MSE of my model and compared them with the model from the practical session.

Practical Session Results:

- Mean squared error (MSE): 0.000001760

- Mean absolute error (MAE): 0.001051122

My MSE & MAE:

- Mean squared error (MSE): 0.000006713
- Mean absolute error (MAE): 0.002116846

```
[154]: model = keras.Sequential([
    keras.layers.LSTM(29, activation='relu', input_shape=(50, 18)),
    keras.layers.Dense(2)
])

[156]: print(model.summary())
Model: "sequential_2"


| Layer (type)    | Output Shape | Param # |
|-----------------|--------------|---------|
| lstm_2 (LSTM)   | (None, 29)   | 5,568   |
| dense_2 (Dense) | (None, 2)    | 60      |


Total params: 5,628 (21.98 KB)
Trainable params: 5,628 (21.98 KB)
Non-trainable params: 0 (0.00 B)
None

[158]: model.compile(optimizer="adam", loss="mse", metrics=["mae"])

[160]: es = EarlyStopping(monitor='val_loss', mode='min', patience=3, verbose=1)
mc = ModelCheckpoint('best_model_LSTM_GOLD.keras', monitor='val_loss', mode='min', verbose=1, save_best_only=True)

[162]: history = model.fit(X_train, y_train, batch_size=20, epochs=10,
                        validation_split=0.1, shuffle=True,
                        verbose=1, callbacks=[es,mc])
Epoch 1/10
1285/1292 - 0s - loss: 0.0237 - mae: 0.0252
```



```
[166]: print("Mean squared error (mse): %.9f %% (scores[0]))")
print("Mean absolute error (mae): %.9f %% (scores[1]))")
Mean squared error (mse): 0.000006713
Mean absolute error (mae): 0.002116846
```



```
[167]: history_dict = history.history
mae_values = history_dict['mae']
val_mae_values = history_dict['val_mae']

epochs = range(1, len(mae_values) + 1)
plt.figure(num=1, figsize=(15,7))
plt.plot(epochs, mae_values, 'b', label='Training Mean Absolute Error(MAE)')
plt.plot(epochs, val_mae_values, marker='o', markeredgecolor='green', markerfacecolor='yellow', label='Validation Mean Absolute Error(MAE)')
plt.xlabel('Epochs', size=18)
plt.ylabel('Mean Absolute Error(MAE)', size=18)
plt.legend()
plt.show()
```

Lab 8

Lab 8 in Week 8, a mock test was conducted with Silver data. For Week 8, students were asked to create templates for the in-class test to be held in Week 9. I created three templates for CNN, LSTM, and MLP techniques.

Results:

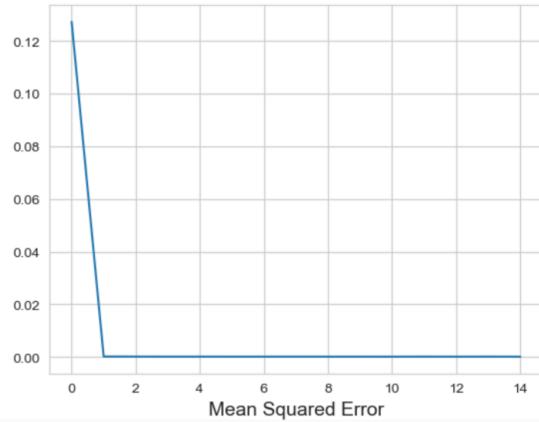
CNN

Plotting the Result Graphs

MSE Training Graphs

```
[107]: #Training MSE (Mean Squared Error)

plt.plot(history.history['loss'])
plt.xlabel('Mean Squared Error', size=14)
plt.show()
```



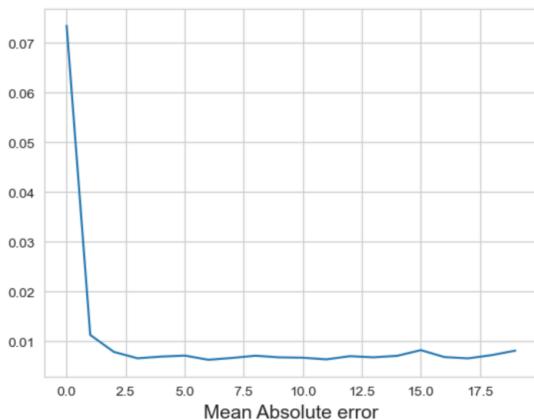
```
Epoch 2/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.0420e-04 - mae: 0.0110 - val_loss: 1.4942e-04 - val_mae: 0.0101
Epoch 3/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.3669e-04 - mae: 0.0118 - val_loss: 2.4954e-04 - val_mae: 0.0134
Epoch 4/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.0764e-04 - mae: 0.0110 - val_loss: 5.1185e-04 - val_mae: 0.0185
Epoch 5/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.9639e-04 - mae: 0.0135 - val_loss: 3.1542e-04 - val_mae: 0.0156
Epoch 6/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.4769e-04 - mae: 0.0123 - val_loss: 4.1447e-04 - val_mae: 0.0174
Epoch 7/15
143/143 ━━━━━━ 0s 2ms/step - loss: 3.4384e-04 - mae: 0.0147 - val_loss: 3.5681e-04 - val_mae: 0.0158
Epoch 8/15
143/143 ━━━━━━ 0s 2ms/step - loss: 4.1869e-04 - mae: 0.0162 - val_loss: 5.1637e-04 - val_mae: 0.0181
Epoch 9/15
143/143 ━━━━━━ 0s 2ms/step - loss: 3.1298e-04 - mae: 0.0140 - val_loss: 1.4145e-04 - val_mae: 0.0093
Epoch 10/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.5210e-04 - mae: 0.0124 - val_loss: 3.5750e-04 - val_mae: 0.0177
Epoch 11/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.1498e-04 - mae: 0.0113 - val_loss: 8.4602e-05 - val_mae: 0.0073
Epoch 12/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.8790e-04 - mae: 0.0129 - val_loss: 7.9444e-04 - val_mae: 0.0264
Epoch 13/15
143/143 ━━━━━━ 0s 2ms/step - loss: 3.1628e-04 - mae: 0.0140 - val_loss: 6.7907e-04 - val_mae: 0.0204
Epoch 14/15
143/143 ━━━━━━ 0s 2ms/step - loss: 2.8178e-04 - mae: 0.0132 - val_loss: 6.3462e-05 - val_mae: 0.0058
Epoch 15/15
143/143 ━━━━━━ 0s 2ms/step - loss: 3.6957e-04 - mae: 0.0152 - val_loss: 7.0861e-05 - val_mae: 0.0062
```

```
[378]: mse, mae = model.evaluate(X_test, y_test, verbose=1)
print("Mean absolute error: %.5f" % mae)
print("Mean squared error: %.5f" % mse)

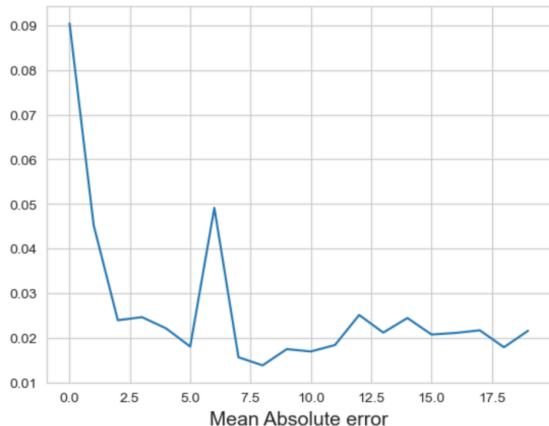
19/19 ━━━━━━ 0s 1ms/step - loss: 1.9928e-04 - mae: 0.0115
Mean absolute error: 0.01154
Mean squared error: 0.00020
```

LSTM

```
[118]: #Training MAE (Mean Absolute Error)
plt.plot(history.history['mae'])
plt.xlabel('Mean Absolute error', size=14)
plt.show()
```



MLP



Lab 9

In the practical session of Week 9, an in-class test was conducted.

Lab 10

Lab 10 in Week 10, Wait-Attention (Flat-Trend) code was developed in the practical session. Students were asked to modify the optimised LSTM model with specific parameters based on their Student ID and include charts showing the training results.

The model must use EarlyStopping() and ModelCheckpoint() callbacks. The required outputs include code screenshots showing callback implementation, model.fit() training output, classification_report(), and the "More detailed Accuracy" graph.

Results:

```
[670]: print(f"Student ID: {SID}")
print(f"Z value (last digit): {Z}")
Student ID: 1919019
Z value (last digit): 9

[671]: # EarlyStopping callback
es = EarlyStopping(
    monitor='val_loss',
    mode='min',
    patience=10,
    verbose=1,
    restore_best_weights=True
)

[672]: # ModelCheckpoint callback
checkpoint_filename = f'best_LSTM_model_SID_{SID}_epochs_{EPOCHS}_batch_{BATCH_SIZE}.keras'
mc = ModelCheckpoint(
    checkpoint_filename,
    monitor='val_loss',
    mode='min',
    verbose=1,
    save_best_only=True
)

[673]: print(f" - EarlyStopping: monitor='val_loss', patience=10")
print(f" - ModelCheckpoint: (checkpoint_filename)")

- EarlyStopping: monitor='val_loss', patience=10
- ModelCheckpoint: best_LSTM_model_SID_1919019_epochs_39_batch_19.keras

[674]: # Train the optimised best LSTM model with modified parameters
history = best_model.fit(
    X_train, y_train,
    batch_size=19, # 19 (Z + 10)
    epochs=39, # 39 (Z + 30)
    validation_split=0.2,
    shuffle=True,
    verbose=1,
    callbacks=[es, mc],
    class_weight=class_weight_dict
)

Epoch 1/39
449/453 - 0s 10ms/step - accuracy: 0.5616 - loss: 0.6858
Epoch 1: val_loss improved from None to 0.68584, saving model to best_LSTM_model_SID_1919019_epochs_39_batch_19.keras
453/453 - 8s 11ms/step - accuracy: 0.5552 - loss: 0.6816 - val_accuracy: 0.6197 - val_loss: 0.6858
Epoch 2/39
451/453 - 0s 9ms/step - accuracy: 0.5482 - loss: 0.6795
Epoch 2: val_loss did not improve from 0.68584
453/453 - 4s 10ms/step - accuracy: 0.5450 - loss: 0.6813 - val_accuracy: 0.4430 - val_loss: 0.7530
Epoch 3/39
452/453 - 0s 10ms/step - accuracy: 0.5847 - loss: 0.6694
Epoch 3: val_loss improved from 0.68584 to 0.65876, saving model to best_LSTM_model_SID_1919019_epochs_39_batch_19.keras
453/453 - 5s 10ms/step - accuracy: 0.5503 - loss: 0.6740 - val_accuracy: 0.6722 - val_loss: 0.6588
Epoch 4/39
448/453 - 0s 10ms/step - accuracy: 0.5487 - loss: 0.6812
Epoch 4: val_loss did not improve from 0.65876
453/453 - 5s 10ms/step - accuracy: 0.5540 - loss: 0.6697 - val_accuracy: 0.5221 - val_loss: 0.7088
Epoch 5/39
452/453 - 0s 12ms/step - accuracy: 0.5076 - loss: 0.6854
Epoch 5: val_loss improved from 0.65876 to 0.43715, saving model to best_LSTM_model_SID_1919019_epochs_39_batch_19.keras
453/453 - 6s 13ms/step - accuracy: 0.5508 - loss: 0.6720 - val_accuracy: 0.8698 - val_loss: 0.4372
Epoch 6/39
448/453 - 0s 10ms/step - accuracy: 0.5685 - loss: 0.6620

Restoring model weights from the end of the best epoch: 5.

[675]: # Load the best saved model
final_model = keras.models.load_model(checkpoint_filename)

[676]: # Evaluate on test data
test_loss, test_accuracy = final_model.evaluate(X_test, y_test, verbose=1)

38/38 - 1s 3ms/step - accuracy: 0.7883 - loss: 0.4805

[677]: # Make predictions
y_pred_prob = final_model.predict(X_test)
y_pred = (y_pred_prob > 0.5).astype(int).flatten()

38/38 - 0s 7ms/step

[678]: print(classification_report(y_test, y_pred,
                                target_names=['Class 0 (Flat)', 'Class 1 (Trend)'],
                                digits=4))

print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")
print(f"Test Error Rate: {(1 - test_accuracy):.4f}")

      precision    recall   f1-score   support
Class 0 (Flat)    0.8874    0.8715    0.8794     1058
Class 1 (Trend)   0.1282    0.1460    0.1365      137
          accuracy                           0.7883    1195
         macro avg    0.5078    0.5087    0.5079    1195
      weighted avg    0.8004    0.7883    0.7942    1195

Test Loss: 0.4805
Test Accuracy: 0.7883
Test Error Rate: 0.2117
```

Lab 11

Lab 12